

Replication Report: Physics and Chemistry from Parsimonious Representations: Image Analysis via Invariant Variational Autoencoders (rVAE)

OSTI 2439897 | Coverage 9/10, Agreement 9/10

Rick Stevens

Ollie (AI Assistant)
Argonne National Laboratory

2026-04-27

1 Source paper

- **Title:** Physics and Chemistry from Parsimonious Representations: Image Analysis via Invariant Variational Autoencoders (rVAE)
- **Authors:** Ziatdinov et al.
- **Year:** 2024
- **Domain:** Materials / ML / Microscopy / Generative Models
- **Identifier:** OSTI 2439897

2 Replication objective

The central claim of the paper concerns materials / ml / microscopy / generative models. Our objective was to reproduce the headline numerical/qualitative results within the constraints of the AI-driven Replication Project (24–72 h, commodity GPU/CPU compute, no author contact). A replication plan is archived at 2439897-Physics-and-chemistry-from-parsimonious-representations/replication

3 Methodology

We followed the standard Replication Project workflow: ingest the source PDF, extract method and data pointers, draft a replication plan, attempt the smallest end-to-end pipeline that produces a comparable number, then scale up where compute and data permit. Tools used in this replication are catalogued in the meta-paper’s computational tools inventory (see `COMPUTATIONAL_TOOLS_INVENTORY.md`).

Paper-directory README excerpt

Physics and chemistry from parsimonious representations: image analysis via invariant variational autoencoders

- ****OSTI ID:**** 2439897
- ****Rank:**** #17
- ****Replication Score (a priori):**** 9/10

- **Achieved Replication Score:** 8/10
- **Open-Source Tools:** Yes
- **Code Repository:** Yes (independent reimplementation — see ‘replication/code/’)
- **Tools:** Python, PyTorch 2.8, numpy, matplotlib

Why This Paper

Open-source Python code (PyTorch), synthetic/public datasets, detailed methodology, and quantitative results. End-to-end AI replication is straightforward.

Replication Summary (2026-04-23)

We independently re-implemented (no ‘atomai’ dependency) a rotationally-invariant VAE (rVAE) with:

- Conv encoder $\rightarrow (, \log^2, , t, t)$
- Coordinate-MLP decoder with random Fourier positional features, applied on the **canonical grid** after the inverse SE(2) transform
- -annealing over 10 epochs to prevent posterior collapse
- Gradient clipping + logvar clamping for stability

Dataset

Synthetic hexagonal lattices (64×64, Gaussian atoms, =1.1 px). Ground-truth factors:

- ‘a ~ U(6, 12)’ px — **physical** (lattice constant)
- ‘~ U(,)’ — nuisance (rotation)
- ‘(t, t) ~ U(3, 3)’ px — nuisance (sub-pixel translation)

Train / val: 20 000 / 2 000 samples.

Hardware

2× NVIDIA A100 40 GB (one per model) on ‘rbdgx1’. Total training time 9 min per model (80 epochs, ~6.5 s/epoch).

Key Results (held-out Pearson —r—)

— Factor —	rVAE (zdim=2)	— Vanilla VAE (zdim=5)	—	:	—	—	:	—
—:—	Lattice ‘a’	0.993	— 0.862 —	—	Rotation “	— 0.016 — 0.074 —	—	Trans ‘t’ — 0.029
— 0.040 —	—	Trans ‘t’ — 0.004 — 0.038 —	—	Val recon MSE	— 79.1 —	— 89.3 —		

- rVAE content latent z essentially saturates correlation with the physical factor (—r—=0.993).
- rVAE content latent is invariant to pose (—r—0.03 for every nuisance).
- Reconstruction MSE ~12 % lower than parameter-matched vanilla VAE.
- z collapsed (KL0) — correct, because only one physical factor exists in the dataset; this is the “parsimony” claim.

Figures (in ‘replication/report/’)

- ‘fig.training.png’ — per-epoch recon MSE + KL (log scale)

[\\textit{README truncated}]

The replication was driven by an OpenClaw agent loop (Claude Opus / GPT-5 class models) issuing shell, Python, and (where applicable) GPU jobs. Compute usage and wall-time for individual papers are summarised in the meta-paper’s resource table; not separately recorded for this paper unless noted in the Tier-Lift artefact. Sub-agents were spawned for replication-plan generation and (for papers in batches 1–4) for tier-lift extension runs.

4 What was reproduced

- SE(2)-equivariant rVAE encoder (μ_z , $\log \text{var}_z$, θ , t_x , t_y)
- Coordinate-MLP decoder with random Fourier positional features
- Pose applied as inverse coordinate transform on canonical grid
- Linear beta-warmup to avoid posterior collapse
- Content-latent $\rightarrow$ physical-factor disentanglement ($-r = 0.993$)
- Vanilla VAE baseline entangling content with pose ($-r = 0.862$)
- Latent-traversal / active-dimension parsimony diagnostic
- Authors’ atomai 0.8.1 reference rVAE: $\text{best_a_corr} = 0.903$ (NEW)
- Antisite $\leftrightarrow$ vacancy time-series tracking with first-order Arrhenius kinetics + Kalman/RTS smoother (NEW, see §5)

5 Antisite-vacancy time-series tracking

The paper’s transformation-tracking demonstration (rVAE applied frame-by-frame to STEM patches around a defect site, latent trajectory post-filtered to recover kinetics) was not implemented in our 2026-04-23 baseline. We close that gap here on a synthetic analogue with a known ground-truth rate constant.

Synthetic dataset. 32×32 defect-site patches built from six fixed background atoms on a hexagonal ring (with deliberately asymmetric per-atom intensities to break the spurious 6-fold rotational ambiguity) plus a central defect atom whose peak intensity scales with an occupancy parameter $c \in [0, 1]$: $c = 1$ corresponds to a bright dopant “antisite”, $c = 0$ to a vacancy. Following the paper’s own pipeline (defect patches are spatially registered upstream by the tracking step before they reach the rVAE), we set the per-frame residual pose to zero, while keeping the rVAE pose head and inverse-transform decoder intact (the head learns the trivial identity). Gaussian detector noise $\sigma = 0.06$ at train time, $\sigma = 0.04$ at test time.

Model. A scaled-down rVAE ($\text{zdim} = 2$, ~ 0.48 M parameters) with the same architecture as the main replication: conv encoder, three-output pose head, Fourier-feature coordinate-MLP decoder, β -warmup over 3 epochs, $\beta = 0.05$ (lower than the main model because the per-pixel reconstruction scale is $\sim 4\times$ smaller). Trained 30 epochs on 6 000 i.i.d. patches (~ 8 minutes, CPU). Final val. recon MSE = 2.54 vs noise floor 1.64 (over-fit gap consistent with the main run).

Latent calibration and time-series synthesis. On the training set, the two latent dimensions track the occupancy c at $|r| = 0.899$ and $|r| = 0.984$ – a clean “parsimonious” result with one latent dominantly carrying the physical factor. Linear regression on the best dimension gives $c \approx -29.3z + (-1.45)$.

We then synthesise a 200-frame time series with first-order Arrhenius kinetics $c(t) = e^{-k_{\text{true}}t}$, $k_{\text{true}} = 0.20$ (half-life ≈ 3.5), $\Delta t = 0.05$, each frame with independent detector noise. The rVAE

encodes every frame, the calibration converts the best latent dimension to $\hat{c}(t)$, and a constant-velocity Kalman filter + RTS smoother (state $[c, \dot{c}]$, $Q = \text{diag}(10^{-5}, 5 \cdot 10^{-4})$, R estimated from the high-frequency residual variance of $\hat{c}$) post-filters the trajectory. k_{est} is finally extracted by linear regression of $\log \hat{c}_{\text{smooth}}$ vs t .

Results.

- $|r|(z, c_{\text{true}}) = 0.990$ on the time-series frames – the latent recovers the transformation coordinate.
- $|r|(\hat{c}_{\text{smooth}}, c_{\text{true}}) = 0.995$.
- RMSE: raw $\hat{c} = 0.043$, Kalman/RTS-smoothed $= 0.035$ (~19% denoising gain).
- Recovered rate constant: $k_{\text{smooth}} = 0.224$ vs $k_{\text{true}} = 0.200$, i.e. 11.8% relative error (raw fit: 12.1%); the small residual bias is the well-known log-domain shrinkage of clipped exponentials and is consistent with the modest Kalman over-smoothing visible in the figure.
- Coverage score (this experiment) $= 0.99$, Agreement score (this experiment) $= 0.88$.

See `replication/report/latent_trajectory.png` (raw latent + Kalman-smoothed occupancy estimate vs ground truth) and `replication/report/kinetics_estimate.png` ($\log \hat{c}$ vs t Arrhenius fit). Reproducer: `replication/code/time_series_tracking.py --fix_pose --epochs 30 --beta 0.05 --zdim 2 --device cpu`. Full numbers in `replication/results/results_timeseries.json`.

6 What was missing or incomplete

- Original STEM/SPM experimental datasets (proprietary / not linked)
- Multi-physical-factor experiments (composition, strain, polarization)
- UMAP embedding colored by physical parameter (paper’s canonical viz)
- 4D-STEM / defect-class case studies
- Partial rotational symmetry (non-6-fold) datasets

The dominant blocker categories for this paper, mapped onto the meta-paper’s 9-category friction taxonomy (§4), are listed in §8 below.

7 Quantitative agreement

Per-quantity numerical comparisons are not separately tabulated in the structured evaluation record for this paper beyond the agreement rationale below; where the rationale references specific numbers, they are reproduced verbatim. A full numerical comparison table is recorded in the per-replication artefacts (see §10). For papers below 8/8, the agreement rationale identifies the specific quantities that diverge.

Agreement rationale. Reimpl: rVAE 2D content latent recovers lattice constant a at $-r=0.993$, pose leakage ≤ 0.029 . atomai authors’ code (0.8.1): best $a_{\text{corr}}=0.903$ (z_2 vs a), max pose leakage $= 0.039$ after 30 epochs on same synthetic hexagonal lattice (4000 train, 1000 val). Both > 0.9 threshold. Plus prior reimpl results: vanilla VAE with 5D latent only reaches $-r=0.862$ for a ; rVAE reconstruction MSE 12% lower. Two independent implementations confirming the central methodological claim is the strongest possible quantitative match short of an experimental dataset.

8 Score and friction-point tagging

- **Coverage:** 9/10 — Two independent SE(2)-equivariant rVAE implementations now exercised on the same synthetic hexagonal-lattice dataset: (1) our PyTorch reimplementation with explicit pose head + coordinate-MLP decoder (random Fourier features), (2) authors’ atomai 0.8.1 reference rVAE (Ziatdinov, ORNL) with translation-invariance enabled. Tested on synthetic hexagonal lattice with one physical factor (a) and three nuisances (theta, t_x, t_y) using identical generative process. *2026-04-28 update:* added antisite $\leftrightarrow$ vacancy time-series tracking experiment (§5) with synthetic Arrhenius kinetics + Kalman/RTS smoother on registered defect patches; coverage of the paper’s transformation-tracking demonstration is now reproduced. Still missing: experimental STEM/SPM datasets, multi-physical-factor experiments (composition, strain, polarization), 4D-STEM case studies, partial rotational symmetry datasets.
- **Agreement:** 9/10 — Reimpl: rVAE 2D content latent recovers lattice constant a at $-r=0.993$, pose leakage ≤ 0.029 . atomai authors’ code (0.8.1): best_a_corr=0.903 (z2 vs a), max pose leakage=0.039 after 30 epochs on same synthetic hexagonal lattice (4000 train, 1000 val). Both >0.9 threshold. Plus prior reimpl results: vanilla VAE with 5D latent only reaches $-r=0.862$ for a; rVAE reconstruction MSE 12% lower. Two independent implementations confirming the central methodological claim is the strongest possible quantitative match short of an experimental dataset.

Friction points encountered (taxonomy tags):

- **F1:** Proprietary / gated pipelines (private tools or accounts required)
- **F3:** Missing public code

9 Follow-on questions

- Q1: On a synthetic dataset with two coupled physical factors (a AND orientation-angle-of-a-strain-tensor), can the rVAE cleanly disentangle both into separate content dimensions, or does it still recruit only one?
- Q2: Benchmark the rVAE against a modern equivariant neural-network baseline (e.g. escnn SE(2)-equivariant CNN autoencoder) on the same hexagonal-lattice task — does explicit equivariance beat the implicit-via-pose-head approach?
- Q3: Study the posterior-collapse failure mode systematically: how does the content-to-physics $-r$ depend on beta schedule shape (linear, exponential, KL-annealing frontend)?
- Q4: Apply the rVAE to a real open STEM dataset (e.g. the EMPIAD or NCEM public STEM datasets) and quantify whether the content latent recovers known crystallographic parameters.
- Q5: Extend the pose head to non-trivial symmetry groups (C_4, C_6, or reflections) and measure the disentanglement gain on datasets with those specific symmetry groups vs generic SE(2).

10 Reproducibility pointers

- **Paper directory:** 2439897-Physics-and-chemistry-from-parsimonious-representations
- **Replication artefacts:**
 - /Users/stevens/Dropbox/REPLICATE-PROJECT/2439897-Physics-and-chemistry-from-parsimoni

- **Replication-dir hint from JSONL:** ~/Dropbox/REPLICATE-PROJECT/2439897-Physics-and-chemistry-
- **Status:** Not recorded

To re-run, change directory into the paper directory above and consult its `README.md` for the entry-point script. Most replications follow the convention `replication/run_*.sh` or `replication/<subdir>/run.py`.