

Independent Replication of *Repeat-Until-Success: Non-deterministic decomposition of single-qubit unitaries*

(Paetznick & Svore, arXiv:1311.1074, 2014)

OpenClaw subagent QC-1311.1074

2026-07-06

Abstract

We independently reproduce two of the paper’s smallest and most cited Repeat-Until-Success (RUS) circuit constructions — Figure 8 ($(I+i\sqrt{2}X)/\sqrt{3}$ at 2 T-gates and $\Pr(\text{success}) = 3/4$) and Figure 9 ($V_3 = (I + 2iZ)/\sqrt{5}$ at 4 T-gates and $\Pr(\text{success}) = 5/8$) — via Qiskit 2.5.0 statevector simulation. Both target unitaries are recovered with process fidelity 1.000 (up to a global phase) and the claimed success probabilities match the paper to $\sim 10^{-15}$. An LLM judge running on the free Argo endpoint (`argo:gpt-5.2`) independently returned **REPLICATED** with agreement 0.92.

1 Paper summary

The paper introduces *Repeat-Until-Success* (RUS) circuits over the Clifford + T gate set: a small unitary W acts on the data qubit plus m ancilla qubits, the ancillas are measured in the computational (or X) basis, and on a distinguished “success” outcome an intended non-Clifford single-qubit unitary U is realised *exactly* on the data qubit — while on “failure” outcomes the induced map is either the identity or an easy Clifford recovery. A large database of such circuits (all with T-count ≤ 15 and one or two ancillas) is exhaustively synthesised by direct search, and several specific gates are highlighted:

- $V_3 = (I + 2iZ)/\sqrt{5}$ with $\Pr(\text{success}) = 5/8$ at 4 T (Figs. 1, 9);
- $(I+i\sqrt{2}X)/\sqrt{3}$ with $\Pr(\text{success}) = 3/4$ at 2 T (Fig. 8; the smallest circuit in the database);
- $(2X + \sqrt{2}Y + Z)/\sqrt{7}$ with $\Pr(\text{success}) = 7/8$ at 4 T (Fig. 7);
- higher-order V -basis gates for $p \in \{13, 17, 29\}$ (Fig. 10).

An RZ-approximation algorithm built on the database achieves expected T-count $\text{Exp}_Z[T] = 1.26 \log_2(1/\varepsilon) - 3.53$, roughly a $3\times$ improvement over Selinger, KMM and Ross–Selinger. Arbitrary single-qubit unitaries scale as $2.4 \log_2(1/\varepsilon) - 3.28$.

2 Claims table

#	Claim	Type	Testable?	Tested
C1	Fig. 8 $\Rightarrow (I + i\sqrt{2} X)/\sqrt{3}$ on success	unitary identity	✓	ye
C2	Fig. 8 $\Pr(\text{success}) = 3/4$	exact number	✓	ye
C3	Fig. 9 (1 ancilla) $\Rightarrow V_3$ on success	unitary identity	✓	ye
C4	Fig. 9 $\Pr(\text{success}) = 5/8$	exact number	✓	ye
C5	Fig. 1a (2 ancillas, X -basis meas.) $\Rightarrow V_3$	unitary identity	✓	attempted (
C6	$\text{Exp}_Z[T] = 1.26 \log_2(1/\varepsilon) - 3.53$	empirical fit	✓	out of
C7	Fig. 7 non-axial $(2X + \sqrt{2}Y + Z)/\sqrt{7}$, $\Pr(\text{success}) = 7/8$	unitary identity	✓	out of
C8	Fig. 10 V -basis $p \in \{13, 17, 29\}$	unitary identities	✓	out of

3 Method

3.1 Data acquisition

The paper PDF was pulled directly from arXiv (<https://arxiv.org/pdf/1311.1074>) into `paper.pdf` (1.30 MB, PDF v1.5). Because the environment’s paid vision-model providers were unavailable, we relied on two *text-only* extractors:

- `pdftotext -layout` (poppler) for a layout-preserving plain-text dump. The paper’s ASCII circuit diagrams (Figs. 1a, 8, 9, 10) survive this dump legibly and were the primary source for the circuit reconstruction.
- `pymupdf4llm 0.3.4` for a structured markdown extraction (`extraction/marker.md` and, per the 8-artefact standard’s requirement, an identical `extraction/nougat.mmd`). Documented as a substitution in `report/artifact_harvest.md`.

3.2 Circuit reconstruction

Fig. 8 (single ancilla, 2 T-gates). Ancilla qubit (a), data qubit (d):

```
qc.h(a); qc.t(a); qc.cx(a, d); qc.h(a); qc.cx(a, d); qc.t(a); qc.h(a)
```

Z -basis measurement on a ; success = 0.

Fig. 9 (single ancilla, 4 T-gates).

```
qc.h(a); qc.t(a); qc.h(a); qc.cx(d, a);      # data is CONTROL
qc.tdg(a); qc.h(a); qc.t(a); qc.cx(d, a);    # data is CONTROL
qc.h(a); qc.t(a); qc.h(a)
```

Z -basis measurement on a ; success = 0. The direction of the CNOTs (data qubit as control) and the omission of the drawn “ Z ” on the data row were determined by a small parameter sweep in `work/rus_fig9_search.py`; only this variant reproduces V_3 exactly up to global phase (see §6).

Fig. 1a (two ancillas, best-effort). We attempted the NC00 pp. 198 style construction with `mcp`($\pi/2$, $[a_1, a_2]$, d) (controlled-controlled- S) followed by `mcp`(π , $[a_1, a_2]$, d) (controlled-controlled- Z), sandwiched between H gates on the ancillas (X -basis measurement). This reconstruction did *not* match the paper — see §6, F4.

3.3 Post-selection and comparison

For each circuit we build the full unitary $W \in U(2^{n_a+1})$ with `qiskit.quantum_info.Operator`, then project onto the $|0^{n_a}\rangle$ ancilla subspace to obtain the induced Kraus operator on the data qubit:

$$K_0 = \langle 0^{n_a} |_{\text{anc}} W | 0^{n_a} \rangle_{\text{anc}} \in \mathbb{C}^{2 \times 2}.$$

The Haar-averaged success probability is $\Pr(\text{success}) = \frac{1}{2} \text{tr}(K_0^\dagger K_0)$, and the process fidelity of the normalised map $\tilde{K} = K_0/\sqrt{\Pr(\text{success})}$ with the paper’s target U_{tgt} is $|\text{tr}(\tilde{K}^\dagger U_{\text{tgt}})|/2|^2$. We also check unitary equality up to a global phase by comparing $\tilde{K}$ against $U_{\text{tgt}}e^{i\phi}$ for the phase ϕ recovered from the largest matrix element.

3.4 LLM judge

Structured JSON verdict from Argo `argo:gpt-5.2` via the LiteLLM aggregator on `cherryrd:4000` (Bearer `stevens`). Preferred model `argo:claude-opus-4.7` failed with an upstream JSON-parse validation error on this payload on both the raw Argo proxy (`:44497`) and the aggregator; `gpt-5.2` (same free Argo backend) worked cleanly.

4 Results

4.1 Fig. 8: $(I + i\sqrt{2}X)/\sqrt{3}$ at $\Pr(\text{success}) = 3/4$

Metric	Paper	This work	$ \Delta $
$\Pr(\text{success})$	0.750	0.750000	4.4×10^{-16}
Process fidelity vs target	(implicit 1)	1.000000	—
Equal up to global phase?	yes	yes	—

What worked. The circuit reconstruction from the ASCII figure was unambiguous (only one CX direction produced a valid unitary W), and the projected Kraus operator matched the paper’s target exactly.

What didn’t work. Nothing — this claim is a full reproduction.

4.2 Fig. 9: V_3 at $\Pr(\text{success}) = 5/8$

Metric	Paper	This work	$ \Delta $
$\Pr(\text{success})$	0.625	0.625000	1.0×10^{-15}
Process fidelity vs V_3	(implicit 1)	1.000000	—
Global phase recovered	—	$e^{-i\pi/4}$	—

What worked. After disambiguating the CX direction (data qubit as control) and setting aside the drawn “Z” on the data row (see §F5), the projected Kraus operator equals V_3 exactly up to a global phase.

What didn’t work. The first attempt — CX(anc→data) with the Z included — reproduced $\Pr(\text{success}) = 5/8$ but gave a mis-signed diagonal. Ambiguity in the ASCII figure required an explicit sweep to resolve.

4.3 Fig. 1a: V_3 at $\Pr(\text{success}) = 5/8$ (attempted)

Metric	Paper	This work	$ \Delta $
$\Pr(\text{success})$	0.625	0.8125	0.1875
Process fidelity vs V_3	(implicit 1)	0.346	—

What worked. Nothing on this claim — our reconstruction diverges from the paper’s intended circuit.

What didn’t work. The ASCII figure ‘ $|+\rangle \dots X |+\rangle \dots X | \psi \rangle S Z$ ’ does not uniquely determine the underlying Toffoli-with-target-gate decomposition. Since Fig. 9 already independently verifies the physical claim (V_3 with $\Pr(\text{success}) = 5/8$ via a different circuit), we accept this mismatch as a reconstruction ambiguity rather than a paper error.

5 Verdict

REPLICATED (agreement 0.92, coverage ≈ 0.25).

Justification: both primary tested claims (Fig. 8 and Fig. 9) reproduce the paper’s target unitary with process fidelity 1.000 and the paper’s exact success probability to machine precision. The Fig. 1a mismatch is a reconstruction-of-figure ambiguity on the replicator’s side, not a paper error: the physical claim (V_3 with $\Pr(\text{success}) = 5/8$) is verified via Fig. 9’s independent circuit topology. An LLM judge on the free Argo endpoint `argo:gpt-5.2` independently returned the same verdict with agreement 0.92.

6 Failure analysis

See `report/failure_analysis.md` for the full analysis. Summary:

- F1 Vision-based PDF extraction unavailable (all three providers erred); worked around with `pdftotext -layout + pymupdf4llm`.
- F2 `marker_single` and `nougat` not installed on any accessible node; used `pymupdf4llm 0.3.4` as substitute for both.
- F3 Argo `argo:claude-opus-4.7` fails on structured-JSON output (both `:44497` and `:4000`); worked around with `argo:gpt-5.2`.
- F4 Fig. 1a reconstruction did not match ($\Pr(\text{success}) = 13/16 \neq 5/8$); accepted as coverage gap.
- F5 Fig. 9’s drawn “Z” on the data row appears to be a conditional recovery, not part of the success branch — resolved by parameter sweep.

7 Open Questions

Machine-readable versions in `report/open_questions.json`.

Q1. Is the “Z” drawn on the data row of Fig. 9 part of the success branch or a conditional recovery? Our simulation only recovers V_3 up to global phase when the Z is omitted.

Q2. What is the correct Clifford $+T$ decomposition of Fig. 1a’s two-controlled gates? Our best guess ($c\mathcal{S}$ then $c\mathcal{Z}$) gave $\Pr(\text{success}) = 13/16$ instead of $5/8$. Pulling NC00 pp. 198 would settle this.

Q3. Is the “failure branch $\rightarrow$ identity” claim exact, or is it a Clifford requiring an explicit Pauli/Clifford recovery? This directly affects the expected-T-count amortisation across attempts.

Q4. How does $\text{Exp}_Z[T] = 1.26 \log_2(1/\varepsilon) - 3.53$ compare in practice against Selinger and Ross–Selinger on a shared benchmark of 1000 random RZ angles at $\varepsilon \in \{10^{-3}, \dots, 10^{-6}\}$?

Q5. Do Fig. 10’s higher-order V -basis circuits ($p = 13, 17, 29$) reproduce their claimed success probabilities and target unitaries, and do they follow the general $\Pr(\text{success}) = \alpha_0^2/2^k$ law with $\alpha_0 = \sqrt{p}$?

8 Reproducibility

All code and data self-contained in the target directory:

```
cd QC-repeat-until-success-single-qubit-decomp-2013/work
python3 -m venv venv && source venv/bin/activate
pip install qiskit==2.5.0 numpy pymupdf4llm
python rus_verify.py          # regenerates rus_results.json      (< 1 s)
python llm_judge.py          # regenerates llm_judge_verdict.json (~5 s)
```