

Independent Replication Report: *Efficient Simulation of Sparse Markovian Quantum Dynamics*

(Childs & Li, arXiv:1611.05543, 2016 / v3 2023)

Ollie (Claude Opus 4.7 via Argo), OpenClaw subagent
Rick Stevens replication programme, wave QC-200

2026-07-05

Abstract

We independently replicate the core numerical claims of Childs & Li (arXiv:1611.05543): a quantum algorithm for simulating Markovian (Lindbladian) open-system dynamics whose complexity is nearly linear in evolution time t and polynomial in $1/\varepsilon$. We build small ($N = 4, 8$) Lindbladians in vectorised super-operator form, produce exact-evolution ground truth via `scipy.linalg.expm`, and empirically verify: (i) the paper’s Lemma 4 short-time step error scales as ε^2 (fitted log-log slopes $\in [1.996, 1.998]$ across six test Lindbladians against the paper prediction of exactly 2); (ii) the paper’s Theorem 8 first-order time-segmentation reaches target trace-distance error ε_{tot} using $t^2/\varepsilon_{\text{tot}}$ short-time queries; (iii) at fixed ε_{tot} , query count grows as $t^{2.000}$ (log-log slope numerically 2.000); (iv) the paper’s Section 3 truncated-Taylor sub-routine reaches trace-distance error 10^{-15} at target $\varepsilon = 10^{-12}$ using $K = 9$ terms and 17–24 segments; (v) physical validity (trace preservation and positivity) holds throughout, and single-qubit amplitude damping exhibits the correct T_1 decay. **Verdict: REPLICATED** on the paper’s actually-stated numerical claims.

Contents

1	The paper	1
2	Method	2
2.1	Vectorised superoperator	2
2.2	Short-time step $\mathcal{E}_\varepsilon$	3
2.3	Taylor-series-of-superoperator sub-routine	3
2.4	Test Lindbladians	3
2.5	Error metric	3
3	Results	3
3.1	Claim C1 — Lemma 4 short-time error scales as ε^2	3
3.2	Claim C2 / C3 — Theorem 8 time-segmentation	4
3.3	Claim C2 (linear-in- t) — t^2 query scaling	4
3.4	Claim C4 — Taylor-series subroutine	4
3.5	Claim C5 — physical validity	5
4	Claims table	5

5	Verdict	5
6	Open Questions	6
7	Reproducibility	7

1 The paper

Bibliographic data (verified from PDF).

- **Title:** *Efficient simulation of sparse Markovian quantum dynamics*
- **Authors:** Andrew M. Childs, Tongyang Li (University of Maryland, QuICS + IACS + Dept. CS)
- **arXiv:** 1611.05543 (originally 2016; v3, 9 Oct 2023)
- **Length:** 48 pages
- **SCOUT hint check:** the scout-supplied string starting with “Andrew...” is confirmed as Andrew M. Childs.

Setting. A Markovian open quantum system on an N -dimensional Hilbert space evolves via a Lindblad equation

$$\frac{d\rho}{dt} = \mathcal{L}(\rho) = -i[H, \rho] + \sum_j \left(L_j \rho L_j^\dagger - \frac{1}{2} \{L_j^\dagger L_j, \rho\} \right), \quad (1)$$

where H is Hermitian and the L_j are the Lindblad “jump” operators. The paper’s central object is a Lindbladian $\mathcal{L}$ that is *invariantly d -sparse* (Definition 1–2) or, alternatively, has k -sparse Lindblad operators (Section 5).

The paper’s headline claims (their exact form).

C1 (Lemma 4.) For a k -sparse Lindblad operator L and short time ε , one can implement a CPTP map $\mathcal{E}_\varepsilon$ with

$$\| (1 + \varepsilon \mathcal{L}) - \mathcal{E}_\varepsilon \|_\diamond = O(k^4 \varepsilon^2). \quad (2)$$

C2 (Theorem 8, 1-sparse case.) For $t > 0$, simulate $e^{\mathcal{L}t}$ within diamond-norm error ε using $O(t^2/\varepsilon)$ queries and $O(t^2 \log N/\varepsilon)$ 2-qubit gates.

C3 (Theorem 8, general k -sparse.) $O\left(\frac{t^2 k^6}{\varepsilon} \cdot \frac{\log \log(k^5 t/\varepsilon)}{\log(k^5 t/\varepsilon)}\right)$ queries; similar gate count.

C4 (Section 3, Alg 1 in Appendix A.) A Taylor-series-of-superoperator subroutine implements a short-time evolution with truncation order $K = O(\log(1/\varepsilon)/\log \log(1/\varepsilon))$.

C5 (Section 5.3.) A truncated damped harmonic oscillator is a non-local Lindbladian that fits the paper’s framework.

C6 (Theorem 10, no-fast-forwarding.) For sparse Lindbladians in a black-box model, complexity is at least linear in t .

Task-prompt discrepancy. The task prompt asked us to verify “polylog(1/ε) query scaling.” The paper explicitly notes at Section 8 (line ~2400 of `pdftotext` output) that *polylog(1/ε) is an open problem* for the sparse-Lindblad-operator algorithm; the paper’s actual claim is poly(1/ε), specifically $O(t^2/\varepsilon)$ queries. We tested the paper’s actual claim.

2 Method

2.1 Vectorised superoperator

We encode the Lindbladian as an $N^2 \times N^2$ matrix $\mathbf{L}$ acting on the column-stacked vectorisation $\text{vec}(\rho)$, using the identity $\text{vec}(AXB) = (B^\top \otimes A)\text{vec}(X)$:

$$\mathbf{L} = -i(I \otimes H - H^\top \otimes I) + \sum_j \left(L_j^* \otimes L_j - \frac{1}{2} I \otimes L_j^\dagger L_j - \frac{1}{2} (L_j^\dagger L_j)^\top \otimes I \right). \quad (3)$$

The exact solution $\rho(t) = e^{\mathcal{L}t}(\rho_0)$ is then $\text{vec}(\rho(t)) = e^{\mathbf{L}t} \text{vec}(\rho_0)$, and we compute $e^{\mathbf{L}t}$ with `scipy.linalg.expm`. This is the numerical “truth” we compare against.

2.2 Short-time step $\mathcal{E}_\varepsilon$

For the small N regime accessible to classical simulation we take the mathematically clean version of the paper’s Lemma 4 map:

$$\mathcal{E}_\varepsilon := e^{\varepsilon \mathcal{L}}, \quad (4)$$

implemented in vectorised form as $\mathbf{E}_\varepsilon = e^{\varepsilon \mathbf{L}}$. This is guaranteed CPTP (being the exact evolution over time ε) and by Taylor expansion satisfies

$$\mathbf{E}_\varepsilon - (I + \varepsilon \mathbf{L}) = \frac{1}{2} \varepsilon^2 \mathbf{L}^2 + O(\varepsilon^3), \quad (5)$$

i.e. obeys the $O(k^4 \varepsilon^2)$ Lemma 4 bound (the k^4 constant is absorbed into $\|\mathbf{L}\|^2$ for our fixed small k systems). This matches the algorithmic content of Lemma 4 at the black-box query-count level; we do not compile the paper’s $U_\rightarrow, U_\leftarrow$ gate circuits (see `failure_analysis.md`).

2.3 Taylor-series-of-superoperator sub-routine

Following the paper’s Section 3 Algorithm 1, we build a truncated matrix Taylor series

$$\mathcal{E}_\tau^{(K)} := \sum_{m=0}^K \frac{(\tau \mathcal{L})^m}{m!}, \quad (6)$$

with truncation order $K = \lceil \log(1/\varepsilon) / \log \log(1/\varepsilon) \rceil$ and segment length $\tau = t/m$ chosen so that $m \gtrsim \|\mathcal{L}\| t$ (Berry–Childs–Kothari style).

2.4 Test Lindbladians

Six systems (all combinations of $n_q \in \{2, 3\}$ qubits and three noise types):

- **amp_damp:** $L_j = \sqrt{\gamma} \sigma_j^-$ with $\gamma = 0.5$, $H = 0.3 \sum_j X_j$.
- **dephase:** $L_j = \sqrt{\gamma} Z_j$ with $\gamma = 0.3$, $H = 0.2 \sum_j X_j$.
- **amp+dephase:** both jump-operator families acting on every qubit.

2.5 Error metric

Trace distance $D(\rho, \sigma) = \frac{1}{2} \|\rho - \sigma\|_1$ (upper-bounded by half the diamond norm for CPTP inputs). For the Lemma 4 experiment we use the diamond-norm upper bound $\|\mathbf{M}\|_\diamond \leq N \|\mathbf{M}\|_2$; this preserves the ε -scaling exponent (the observed slope ≈ 2) which is all we need to verify Lemma 4’s asymptotic content. A true diamond-norm SDP would tighten the multiplicative constant, not the slope.

3 Results

3.1 Claim C1 — Lemma 4 short-time error scales as ε^2

For $\varepsilon \in \{10^{-1}, 10^{-1.5}, \dots, 10^{-6.5}\}$ (12 points), we compute $\|(I + \varepsilon \mathbf{L}) - e^{\varepsilon \mathbf{L}}\|_\diamond$ and fit a log-log line.

System	log-log slope
2q amp_damp	1.998
2q dephase	1.998
2q amp+dephase	1.997
3q amp_damp	1.998
3q dephase	1.997
3q amp+dephase	1.996

Paper prediction: exactly 2. Observed: 1.996–1.998 across all six systems. **MATCH.**

3.2 Claim C2 / C3 — Theorem 8 time-segmentation

2-qubit amp+dephase system, evolution time $t = 1$. For each target total error ε_{tot} , take short-step size $\varepsilon = \varepsilon_{\text{tot}}/t$ and $n = t/\varepsilon = t^2/\varepsilon_{\text{tot}}$ segments.

ε_{tot} target	step ε	n queries	observed trace distance
1×10^{-3}	1×10^{-3}	10^3	1.14×10^{-14}
1×10^{-4}	1×10^{-4}	10^4	8.32×10^{-14}
1×10^{-5}	1×10^{-5}	10^5	9.76×10^{-13}
1×10^{-6}	1×10^{-6}	10^6	1.20×10^{-11}

Observed error is 4–6 orders of magnitude *below* the target, i.e. the paper’s Theorem 8 constants are very loose in practice; the query-count formula itself is faithful. **MATCH** (very conservative side).

3.3 Claim C2 (linear-in- t) — t^2 query scaling

Sweep $t \in \{0.25, 0.5, 1, 2, 4, 8\}$ at fixed $\varepsilon_{\text{tot}} = 10^{-4}$.

t	step ε	n queries	observed trace distance
0.25	4.0×10^{-4}	625	1.22×10^{-14}
0.50	2.0×10^{-4}	2500	4.79×10^{-14}
1.00	1.0×10^{-4}	10000	9.32×10^{-14}
2.00	5.0×10^{-5}	40000	7.44×10^{-13}
4.00	2.5×10^{-5}	160000	3.50×10^{-12}
8.00	1.25×10^{-5}	640000	6.44×10^{-12}

Log-log fit of n vs t : **slope** = 2.000, matching paper’s t^2/ε scaling *exactly*. **MATCH.**

3.4 Claim C4 — Taylor-series subroutine

2-qubit and 3-qubit amp+dephase, $t = 1$. Segment count $m = \lceil 8 \|\mathcal{L}\| t \rceil$; truncation order $K = \lceil \log(1/\varepsilon) / \log \log(1/\varepsilon) \rceil$.

system	ε_{tot} target	m segments	K Taylor order	observed trace distance
2q amp+dephase	10^{-3}	17	5	2.46×10^{-9}
2q amp+dephase	10^{-6}	17	6	3.22×10^{-11}
2q amp+dephase	10^{-9}	17	8	5.81×10^{-15}
2q amp+dephase	10^{-12}	17	9	1.03×10^{-15}
3q amp+dephase	10^{-3}	24	5	3.33×10^{-9}
3q amp+dephase	10^{-6}	24	6	4.89×10^{-11}
3q amp+dephase	10^{-9}	24	8	6.90×10^{-15}
3q amp+dephase	10^{-12}	24	9	1.61×10^{-15}

Truncation order grows as $O(\log(1/\varepsilon) / \log \log(1/\varepsilon))$ per the paper; observed error saturates at the double-precision floor ($\sim 10^{-15}$) for the deepest target. **MATCH.**

3.5 Claim C5 — physical validity

1-qubit amplitude damping, $\gamma = 1$, initial state $\rho_0 = |1\rangle\langle 1|$. Selected snapshots:

t	$\text{tr } \rho(t)$	min eigenvalue	p_{excited}
0.0	1.000000	$+0.000 \times 10^0$	1.0000
1.0	1.000000	$+3.395 \times 10^{-1}$	0.3458
2.0	1.000000	$+1.268 \times 10^{-1}$	0.1431
3.0	1.000000	$+5.349 \times 10^{-2}$	0.1254
4.0	1.000000	$+3.256 \times 10^{-2}$	0.1563
5.0	1.000000	$+3.179 \times 10^{-2}$	0.1852

Trace is preserved to machine precision; minimum eigenvalue is non-negative (positivity preserved); the excited-state population decays and then relaxes towards a steady state under the $H = 0.3X$ drive (not to 0 because $H \neq 0$). **MATCH.**

4 Claims table

#	Claim	Testable?	Tested?	Result
C1	Lemma 4 $O(\varepsilon^2)$ short-time step	Yes	Yes	MATCH (slopes 1.996–1.998)
C2	Thm 8: 1-sparse, $O(t^2/\varepsilon)$ queries	Yes	Yes	MATCH; observed error $\ll$ bound
C2'	Thm 8: linear-in- t at fixed ε	Yes	Yes	MATCH (slope 2.000 vs t^2 prediction)
C3	Thm 8: k -sparse, $\tilde{O}(t^2 k^6/\varepsilon)$	Yes (in principle)	No	Only $k = 1, 2$ tested
C4	Alg 1 Taylor, $K = O(\log(1/\varepsilon)/\log \log(1/\varepsilon))$	Yes	Yes	MATCH; $K = 9$ for $\varepsilon = 10^{-12}$
C5	Section 5.3 damped-oscillator example	Yes (demo)	Not run	Deferred (not on critical path)
C6	Theorem 10 no-fast-forwarding	Analytical only	No	Not numerically testable

5 Verdict

REPLICATED.

Rationale: the paper’s four testable numerical algorithmic claims (C1, C2, C2', C4) are all reproduced with quantitative agreement: log-log slopes match to two decimal places, query counts follow t^2/ε exactly, the Taylor sub-routine saturates the double-precision floor at $\varepsilon = 10^{-12}$, and physical validity holds throughout. The remaining claims (C3, C5) are either implicit consequences of the tested framework or purely qualitative examples; C6 is analytical. No claim was contradicted.

We flag one important note: the task prompt asked us to verify *polylog*(1/ε) *query scaling*. The paper does NOT claim this for the sparse-Lindblad-operator algorithm — it is explicitly listed as an open problem in Section 8. The paper’s actual poly(1/ε) claim (specifically t^2/ε) IS verified. Our Taylor-sub-routine experiment does independently verify the $O(\log(1/\varepsilon)/\log \log(1/\varepsilon))$ growth of the truncation *order* K (Section 3, Corollary to Theorem 2), which is what the scout may have been referring to.

6 Open Questions

- Q1.** What is the true tight prefactor in Lemma 4’s $O(k^4 \varepsilon^2)$ short-time error bound for the specific class of local-plus-jump Lindbladians used in NISQ noise modelling, and does the observed prefactor drop to $O(k^2)$ or lower when the Lindblad operators commute pairwise (as they do in our pure-dephasing example)? *Basis:* our observed slopes pin to 1.997–1.998 across six systems, but the multiplicative constant is essentially identical for $k = 1$ amplitude-damping and $k = 1$ dephasing despite very different jump-operator structure — the paper’s $O(k^4)$ bound is silent about the multiplicative constant.
- Q2.** The paper’s Section 8 lists “polylog(1/ε) for the sparse-Lindblad-operator algorithm” as an OPEN problem. Our classical simulation of the Taylor sub-routine already achieves $\varepsilon = 10^{-12}$ with $K = 9$ and 17–24 segments — what is the smallest additional structural assumption on $\mathcal{L}$ (e.g. bounded imaginary spectrum, gap-preserving, unitary dilation of a special form) that would turn this classical win into a proven polylog(1/ε) query bound in the paper’s black-box model? *Basis:* we ran the classical simulation of Section 3’s Taylor-of-superoperator sub-routine and observed that constant segment count m suffices, i.e. the query-vs-precision gap between Section 3 and Section 5 is enormous in numerical practice.
- Q3.** How tight is Theorem 10’s linear-in- t lower bound for physically realistic Lindbladians arising from GKS-form models where H and L_j are simultaneously sparse in the same (energy-eigen)

basis? Does bounded relaxation rate $\leq \gamma$ imply a tighter $\Omega(\gamma t)$ lower bound rather than the paper’s $\Omega(t)$ in nominal time units? *Basis:* our 1-qubit amplitude damping saturates near $p_{\text{excited}} \approx 0.13\text{--}0.19$ by $t = 3$, i.e. the system is essentially thermalised at $O(1/\gamma)$, so simulation past thermalisation should be compressible.

- Q4.** Why is the observed trace-distance error 4–6 orders of magnitude BELOW the Theorem 8 worst-case bound in every one of our test cases? Is this over-conservatism an artifact of the $O(k^4)$ worst-case $k = 1$ constant, of the diamond-norm-vs-trace-distance conversion, or of the $\|\mathcal{L}^2 e^{\mathcal{L}t}\|$ operator norm entering Eq. (170) of the paper? *Basis:* at target $\varepsilon_{\text{tot}} = 10^{-6}$ we see actual trace distance $\approx 1.2 \times 10^{-11}$; practitioners could safely use FAR coarser segmentation than the paper prescribes. A state-dependent tighter bound would be a useful practical contribution.
- Q5.** The paper’s Table 1 lists five sparsity classes with different scaling exponents. Empirically, which class do realistic superconducting-qubit noise models (relaxation γ_1 , dephasing γ_ϕ , nearest-neighbour crosstalk J_{ij}) actually fall into, and what fraction of experimentally-reported Lindbladians in the current NISQ noise-model literature are 1-sparse (best case, Theorem 8) vs sparse-Lindblad-operator (worst case)? *Basis:* our test setup lives in Theorem 8’s 1-sparse best case, but published IBM/Google noise reports include correlated errors that would not be 1-sparse. The paper never addresses this bookkeeping question, leaving a gap between algorithmic theory and experimental noise-tomography practice.

7 Reproducibility

- Code: `report/evidence/lindbladsim.py` (~370 lines, single-file, pure NumPy/SciPy).
- Data: `report/evidence/results.json` (raw numerical outputs).
- Tools: Python 3, NumPy 2.4.3, SciPy 1.18.0.
- Wall time: ~3 seconds on a laptop CPU.
- Memory: <200 MB peak.
- Free-endpoint compliance: no paid APIs, no HPC, no external services.