

Independent Replication Report:
Quantum Algorithms for some Hidden Shift Problems
(van Dam, Hallgren, Ip; arXiv:quant-ph/0211140)

QC-200 Replication (independent, numpy exact-statevector)

2026-07-05

1 Paper summary

van Dam, Hallgren and Ip [?] give the first quantum algorithms for the *hidden shift problem*: given oracle access to $f(x) = g(x + s)$ where g is a completely known function on an abelian group G and $s \in G$ is an unknown shift, find s . Three families are studied:

- (A) g is a nontrivial multiplicative character of a finite field $\mathbb{F}_q$ (the flagship case is the shifted Legendre symbol on $\mathbb{F}_p$). A single-query algorithm succeeds with probability $(1 - 1/q)^2$.
- (B) g is a multiplicative character of $\mathbb{Z}/n\mathbb{Z}$ (Jacobi-symbol analogue), where the shift need not be unique.
- (C) Same as (B), but with n unknown.

Their central circuit (Fig. 1 of the paper) is

$$|0\rangle \xrightarrow{F_G} |x\rangle \mapsto f(x)|x\rangle \xrightarrow{F_G} |x\rangle \mapsto \hat{g}^{-1}(x)|x\rangle \xrightarrow{F_G} \xrightarrow{\text{measure}} -s,$$

where F_G is the QFT on G and $\hat{g}$ is the (classically computable) Fourier transform of the known function g .

2 Claims table

ID	Claim	Testable?	Tested?
C1	Algorithm 1 (Sec. 4) solves the shifted-multiplicative-character problem over $\mathbb{F}_q$ with success probability $(1 - 1/q)^2$ from ONE oracle query.	Yes (small p)	Yes
C2	Specialisation to Legendre symbol on $\mathbb{F}_p$: recover s in polynomial time.	Yes ($p = 13$)	Yes
C3	The same circuit template (QFT-oracle-QFT with phase-uncompute $\hat{g}^{-1}$) recovers s on any abelian G provided $\hat{g}$ is nowhere zero.	Yes ($G = \mathbb{Z}_N$, $N = 8, 16, 32$)	Yes
C4	Exponential quantum-vs-classical separation in oracle complexity: $O(1)$ vs $\Omega(\log N)$ classical queries to recover s information-theoretically.	Yes (info-theoretic + empirical distinguisher)	Yes
C5	Extension to $\mathbb{Z}/n\mathbb{Z}$ with possibly non-unique shifts (Sec. 5).	Yes but scope-cap	Partial (skip)
C6	Extension to unknown n (Sec. 5.2) via approximate Fourier sampling.	Yes but scope-cap	Not tested
C7	Hidden coset framework (Sec. 6) unifying HSP + hidden shift.	Definitional (no numerical value)	Not applicable

Testing scope: this replication targets the core testable numerical claims C1–C4. The framework-level C5–C7 are algorithmic generalisations without a single numerical headline number.

3 Method

3.1 Environment

- Host: CherryRd, macOS 25.3.0, Python 3.13.
- Venv: `~/openclaw/workspace/venvs/qc-replicate/`.
- Packages: `numpy 2.5.1`, `sympy` (unused at runtime; imported for future extension).
- No Qiskit / Cirq were needed: the paper’s circuits act on ≤ 32 amplitudes, so an exact statevector implemented as $N \times N$ numpy matrices is faster and truer to the paper than a gate-decomposed simulator.
- No GPU / HPC.

3.2 Marker / Nougat extractions

Neither `marker-pdf` nor `nougat-ocr` was installed on the host, and their pip installs pulled `torch` which the local wheel index could not resolve within the wave-time budget (see `failure.analysis.md`). `extraction/marker.md` and `extraction/nougat.mmd` were therefore produced from `pdftotext` of the source PDF with a documented disclaimer header, per the “fallback to pdftotext with disclosure” pattern used by other QC-200 backfills. The raw `pdftotext` outputs (`work/paper.txt` and `work/paper_layout.txt`) are retained as ground-truth backups.

3.3 Numerical replication

Code: `report/evidence/hidden_shift_zn.py` (~300 LOC). Runs in ~ 1 s on a laptop.

Algorithm 1 on $\mathbb{Z}_N$

For each $N \in \{8, 16, 32\}$ we choose a known function $g : \mathbb{Z}_N \rightarrow \mathbb{C}$ and a random hidden shift s , then execute the four steps of the paper’s Fig. 1 circuit as exact matrix–vector multiplications:

1. Amplitude-encode $|\psi\rangle = N^{-1/2} \sum_x g(x+s)|x\rangle$.
2. Apply the exact N -dim DFT matrix F .
3. Apply the diagonal phase $\text{diag}(\overline{\hat{g}(y)}/|\hat{g}(y)|)$.
4. Apply $F^\dagger$ and read the argmax of $|\psi|^2$.

Two families of g were tested per N :

Chirp g : $g(x) = \exp(2\pi i a x^2 / (2N))$ (Zadoff-Chu sequence) with a tuned so $|\hat{g}| \equiv 1$. This is the cleanest analogue of the paper’s multiplicative character: *flat* spectrum and known closed form. Predicted outcome: deterministic $(N-s) \bmod N$ with probability 1.

Boolean g : random $g : \mathbb{Z}_N \rightarrow \{\pm 1\}$ rejection-sampled until $|\hat{g}(y)| > 0$ for all y . This is the task-brief’s request for a Boolean g and lets us stress-test the algorithm on functions with non-flat spectrum. Predicted: s recoverable as argmax but with peak probability < 1 .

Ten independent random shifts were run per $(N, g\text{-type})$ pair.

Shifted Legendre symbol on $\mathbb{F}_{13}$

Exact p -dim statevector implementation of Algorithm 1 with $\chi = (\frac{\cdot}{13})$. Ten random nonzero shifts.

Classical query lower bound

For each N we compute both the information-theoretic $\lceil \log_2 N \rceil$ bound (each classical query into a ± 1 oracle yields ≤ 1 bit) and an empirical average – over 200 random non-adaptive size- k query sets – of the worst-case number of shifts still consistent with the observed answer vector.

4 Results vs. paper

Instance	Trials	Successes	Empirical peak prob.	Paper prediction	Verdict
$\mathbb{Z}_8$ chirp g	10	10/10	1.000	1 (flat spec.)	MATCH
$\mathbb{Z}_{16}$ chirp g	10	10/10	1.000	1 (flat spec.)	MATCH
$\mathbb{Z}_{32}$ chirp g	10	10/10	1.000	1 (flat spec.)	MATCH
$\mathbb{Z}_8$ Boolean g	10	10/10 (argmax)	0.729	— (not in paper)	MATCH scope
$\mathbb{Z}_{16}$ Boolean g	10	10/10 (argmax)	0.699	—	MATCH scope
$\mathbb{Z}_{32}$ Boolean g	10	10/10 (argmax)	0.804	—	MATCH scope
Legendre $\mathbb{F}_{13}$	10	10/10 (argmax)	0.923	$(1 - 1/13)^2 = 0.852$	MATCH

The Legendre empirical peak probability 0.923 exceeds the paper’s $(1 - 1/13)^2 = 0.852$. This is not a discrepancy: the paper’s number bounds the overall success probability of the ENTIRE algorithm including Lemma 1 (which fails with probability $1/p$ due to $\chi(0) = 0$). Our exact-statevector implementation skips the Lemma 1 measurement branch and directly builds the phase-encoded state, so we see the conditional success probability given the amplitude construction succeeded. Multiplying by $(1 - 1/p)$ for the Lemma-1 success probability gives $0.923 \cdot (1 - 1/13) = 0.852$, exactly matching the paper.

N	Quantum oracle queries	Classical info-theoretic bound	Speedup
8	2 (Lemma 1)	$\lceil \log_2 8 \rceil = 3$	exponential in $\log N$
16	2	$\lceil \log_2 16 \rceil = 4$	exponential in $\log N$
32	2	$\lceil \log_2 32 \rceil = 5$	exponential in $\log N$

For the Legendre variant the classical separation is much larger: the best known classical algorithm for the shifted Legendre symbol is subexponential [?]; the quantum algorithm is polynomial-time.

5 Per-claim what worked / what did not work

C1 (finite-field Alg. 1 success probability). Worked cleanly on $\mathbb{F}_{13}$. Empirical $p_{\max} = 0.923$; theory $(1 - 1/13)^2 = 0.852$; discrepancy fully explained by our Lemma-1-conditioning (see Results section). Would need Chinese-Remainder plumbing to test on prime powers $q = p^r$, which we did not do.

C2 (Legendre recovery). Worked. 10/10 random nonzero shifts on $\mathbb{F}_{13}$ recovered. The measurement outcome $(p - s) \bmod p$ matched the theory prediction in all trials.

C3 (generalisation to $\mathbb{Z}_N$). Worked for three group sizes $N = 8, 16, 32$. With the chirp g (flat-spectrum, mimics a multiplicative character), recovery is deterministic (peak prob. = 1). With Boolean g (task-brief-requested variant, non-flat but nowhere-zero spectrum), recovery is probabilistic (≈ 0.7 – 0.8) but the correct shift is always the argmax over 10 trials.

C4 (query separation). Confirmed both info-theoretically ($k \geq \lceil \log_2 N \rceil$ classical queries required) and empirically: at $k = 1$ – 2 classical non-adaptive queries the worst-case ambiguity class contains many shifts, only shrinking to a single shift once k reaches $\log_2 N$.

C5–C6 (Sec. 5 extensions to $\mathbb{Z}/n\mathbb{Z}$ composite and unknown- n). *Not tested.* The Jacobi-symbol case requires Chinese-Remainder decomposition of the character and its Fourier transform (paper’s Sec. 5.1); the unknown- n case additionally requires the approximate Fourier sampling machinery of [?] (paper’s Sec. 2.3). Both extensions are algorithmically documented in the paper but their concrete implementation was outside the wave-time budget. Recorded as a follow-up.

C7 (hidden coset framework). Definitional; nothing to numerically test.

6 Verdict

REPLICATED. The paper’s core algorithm reproduces on three group sizes ($\mathbb{Z}_8, \mathbb{Z}_{16}, \mathbb{Z}_{32}$) with polylog quantum query count, and the flagship Legendre-symbol specialisation recovers the shift on $\mathbb{F}_{13}$ with the paper-predicted success probability. All numerical predictions we could test matched, once the Lemma-1 conditioning is accounted for. The Sec. 5 extensions (composite n , unknown n) were not implemented; the definitional Sec. 6 (hidden coset) has no numerical content.

Open Questions

Q1. The Lemma-1 amplitude-construction success probability $(1 - 1/q)$ is the ONLY probabilistic branch in the paper’s finite-field algorithm; conditioned on Lemma 1 succeeding, our simulation shows $p_{\max} = (1 - 1/q)$ additional probability of the correct measurement, hitting the paper’s $(1 - 1/q)^2$ product. Is there a one-query variant of Lemma 1 (say, using amplitude amplification across the two calls it makes) that removes the $1/q$ Lemma-1 penalty and lifts the total success probability to $\geq 1 - 1/q$?

Q2. For Boolean g on $\mathbb{Z}_N$ we observed peak probabilities ~ 0.7 – 0.8 instead of the flat-spectrum 1. Which combinatorial property of g tightly controls the peak probability – is it $\min_y |\hat{g}(y)|^2$, the flatness ratio $\max/\min$, or some higher moment? A closed-form lower bound in terms of a computable invariant of g would be useful for choosing oracles in cryptographic reductions.

Q3. We used the Zadoff-Chu chirp $g(x) = \exp(2\pi i ax^2/(2N))$ as a flat-spectrum stand-in for a multiplicative character on $\mathbb{Z}_N$ (which is not a field). Does the hidden-shift algorithm using such a chirp g inherit the paper’s Legendre-symbol cryptographic implications (attacking algebraically homomorphic schemes), or does the field structure of $\mathbb{F}_p$ do essential work beyond providing a flat-spectrum g ?

Q4. The paper reduces the shifted-Legendre problem to the dihedral HSP under Conjecture 2.1 of [?]. Our exact simulation on $p = 13$ solves the shift problem directly (no dihedral reduction). Can the empirical peak-probability curve from small- p brute-force runs be used to test/tighten Conjecture 2.1 numerically? E.g. do the residual amplitude patterns predicted by the conjecture appear in the $|\psi|^2$ distribution at $p = 11, 13, 17, 19$?

Q5. The Sec. 5.2 unknown- n variant requires approximate Fourier sampling with parameters $m = \Omega(n^2/\epsilon^2)$, $q = \Omega(m/\epsilon)$. For which (ϵ, n) regimes does this approximation degrade the shift-recovery probability below $1/\text{poly}(\log n)$, i.e. what is the smallest polylogarithmic sample size that still recovers s with constant success probability in the unknown- n case? This would clarify the practical scaling of the “unknown n ” extension, which we did not implement.

References

- [1] W. van Dam, S. Hallgren, L. Ip. *Quantum Algorithms for some Hidden Shift Problems*. arXiv:quant-ph/0211140 (2002).
- [2] D. Boneh and R. Lipton (as cited by vDHI02) – subexponential classical algorithms for algebraically homomorphic cryptosystems.
- [3] L. Hales and S. Hallgren, “An improved quantum Fourier transform algorithm and applications,” FOCS 2000.