

Independent Replication of “Quantum algorithm for a generalized hidden shift problem”

Childs & van Dam, arXiv:quant-ph/0507190v1 (2005)

QC-200 replication (subagent under OpenClaw)

2026-07-05

Verdict

REPLICATED. The paper’s two central quantitative claims are reproduced numerically end-to-end. Specifically:

- The closed-form pretty-good-measurement (PGM) success probability (Eq. 15 of the paper) is computed exactly for $N \leq 64$ by enumerating all η_w^x ; the success probability is bounded away from zero when $M = \lfloor N^{1/k} \rfloor$ with $k \geq 3$ (Lemma 2 regime), and degrades if M is held below that threshold, matching the paper’s theoretical assertions.
- The PGM POVM built numerically as $E_j = \Sigma^{-1/2} \sigma_j \Sigma^{-1/2}$ gives an operational success probability $\text{Tr}(E_s \sigma_s)$ that matches the analytic Eq. (15) to machine precision ($|\text{diff}| < 7 \times 10^{-16}$) for all small cases we tested.
- An end-to-end Qiskit statevector simulation of the Naimark-dilated PGM circuit on 8–10 qubits reproduces the analytic success probability within Monte-Carlo error ($|\text{emp} - \text{Eq. 15}| \leq 0.043$ at ≤ 500 shots for $(N, M, k) \in \{(3, 2, 2), (4, 2, 2), (3, 2, 3)\}$).

See Section ?? for the full numerical comparison.

1 Paper summary

Childs and van Dam introduce the *generalized hidden shift problem* (GHSP): given $f : \{0, \dots, M-1\} \times \mathbb{Z}_N \rightarrow S$ with $f(b, x) = f(b+1, x+s)$ for $b = 0, \dots, M-2$ and $f(b, \cdot)$ injective for each b , recover the hidden shift $s \in \mathbb{Z}_N$. The problem interpolates between the dihedral hidden subgroup problem ($M = 2$, believed hard) and the abelian HSP ($M = N$, solved by Shor-style Fourier sampling).

The paper’s algorithm has two coupled pieces:

1. **Quantum:** implement the pretty good measurement (PGM) on k copies of the coset-state ensemble ρ_s (Eqs. 2–3). The PGM has success probability given exactly by Eq. (15).
2. **Classical:** inside the PGM, reduce the required “quantum sampling” step to an instance of k -dimensional integer programming (the “matrix sum problem”); use Lenstra’s algorithm to solve it in $2^{O(k^3)}$ time.

Together with Lemma 1 (upper/lower bounds on PGM success in terms of matrix-sum solvability) and Lemma 2 (constant lower bound on $\Pr(1 \leq \eta_w^x \leq 4)$ when $M = \lfloor N^{1/k} \rfloor$), this yields Theorem 3: GHSP is solvable in $\text{poly}(\log N)$ time on a quantum computer for any fixed $\varepsilon > 0$ when $M \geq N^\varepsilon$.

2 Claims table

ID	Claim	Testable?	Tested in this replication?
C1	Eq. (15): PGM success probability closed form $\Pr(\text{success}) = \frac{1}{M^k N^{k+1}} \sum_x (\sum_w \sqrt{\eta_w^x})^2$	Yes, exactly, by enumeration for small N	Yes: computed analytically <i>and</i> operationally as $\text{Tr}(E_s \sigma_s)$; equal to $\sim 10^{-16}$ (Sec. ??, Test B).
C2	Lemma 2: $\Pr(1 \leq \eta_w^x \leq 4)$ is bounded below by a constant for $M = \lfloor N^{1/k} \rfloor$, $k \geq 3$, large N	Yes, finite- N trend	Yes: for $k = 3, N = 27 \Rightarrow \Pr = 0.651$; $N = 32 \Rightarrow 0.558$. For $k = 4$ up to $N = 24$, $\Pr \geq 0.489$ (Test A).
C3	Theorem 3: GHSP solvable in $\text{poly}(\log N)$ time when $M \geq N^\varepsilon$	Complexity claim; <i>small-N regime cannot show asymptotics</i> , but success prob $\geq$ constant is verifiable	Partial: success-prob part verified operationally at small N ($k = 3, N = 32$: $p = 0.547$; $k = 4, N = 24$: $p = 0.479$). Lenstra IP subroutine NOT implemented (would require external IP solver); we brute-force enumerate matrix-sum solutions instead.
C4	The PGM is optimal for this ensemble (paper cites [2,3])	Yes, by comparing to numerical optimum	Not tested (out of scope: would require solving the Yuen–Kennedy–Lax SDP; interesting follow-up).
C5	For $M \ll N$, abelian Fourier sampling followed by comp.-basis measurement gives success M/N (paper’s paragraph after Eq. 5)	Yes, one-copy Fourier-sampling probability check	Not tested (single-copy negative result; the paper’s positive claim is the multi-copy PGM).
C6	The generalized problem on $\{0, \dots, M - 1\}$ reduces to that on any smaller $\{0, \dots, M' - 1\}$ (proof of Thm. 3)	Yes, structural	Verified conceptually: our code parameterizes over arbitrary M ; the reduction is just “ignore the top values of b ”, which is trivially correct.

3 Method (numbered, reproducible)

3.1 Environment

- Host: CherryRd, macOS Darwin 25.3.0
- Python 3.14.6, virtualenv at `.venv/`
- Packages: qiskit 2.5.0, qiskit-aer, numpy 2.5.1, scipy, matplotlib
- Zero paid-API/inference use; simulation is entirely classical numpy + Qiskit AerSimulator (statevector).

3.2 Steps executed

1. Fetch the paper: `curl -sL https://arxiv.org/pdf/quant-ph/0507190 -o paper.pdf`
2. `pdftotext paper.pdf work/paper.txt` (used to derive extraction/marker.md and extraction/nougat.mmd; note that neither Marker nor Nougat is installed on this host, and no pre-parsed copy of quant-ph/0507190 exists in the local corpus, so these two extractions are honest pdftotext-derived fallbacks and are labeled as such in their headers).
3. Skim + extract algorithm: identified Eq. (15) as the single most-checkable quantitative claim (closed-form PGM success probability), and Lemma 2 as the paper’s key combinatorial ingredient.
4. Implement `childs_vandam_pgm.py` (in `code/`), consisting of three tests:
 - **Test A — Sweep.** Enumerate η_w^x for all $(x, w) \in \mathbb{Z}_N^k \times \mathbb{Z}_N$ via a batched Fourier-domain convolution (per-coordinate count vectors combined coordinate-wise). Compute Eq. (15) and $\Pr(1 \leq \eta_w^x \leq 4)$. Sweep $k \in \{2, 3, 4\}$, $N \in \{4, \dots, 64\}$ with $M = \lfloor N^{1/k} \rfloor$.
 - **Test B — Operational check.** For small (N, M, k) , build $\rho_s^{\otimes k}$ explicitly in a $(MN)^k$ -dimensional Hilbert space (numpy density matrices); construct the PGM POVM $E_j = \Sigma^{-1/2} \sigma_j \Sigma^{-1/2}$ (pseudo-inverse-sqrt over the ensemble support); compute $\text{Tr}(E_s \sigma_s)$ and compare to the analytic Eq. (15).
 - **Test C — Qiskit end-to-end.** For $(N, M, k) \in \{(3, 2, 2), (4, 2, 2), (3, 2, 3)\}$: for each shot, sample $x_1, \dots, x_k \sim \text{Unif}(\mathbb{Z}_N)$ (this is the classical marginal inside $\rho_s^{\otimes k}$), prepare $|\Phi_{x,s}\rangle = \bigotimes_i |\phi_{x_i,s}\rangle$ on a Qiskit `QuantumCircuit`, apply the Naimark-dilated PGM unitary U on system+ancilla, measure the ancilla, sample outcome j . Aggregate $\Pr(j = s)$ over shots.
5. Save all outputs to `report/evidence/results.json` and `report/evidence/run_log.txt`; produce plots via `code/make_figures.py`.

3.3 Naimark dilation implementation (Test C)

The PGM POVM elements E_j on this ensemble are generically full-rank (rank up to N^k , block-diagonal per x -block per paper Eq. (21)), so a naive “ $E_j = |f_j\rangle\langle f_j|$ ” rank-1 factorization is *wrong* for this problem. We use the correct Kraus construction:

$$K_j = \sqrt{E_j}, \quad \sum_j K_j^\dagger K_j = P_\Sigma \text{ (projector on ensemble support)}$$

and build the Naimark unitary U on system $\otimes$ ancilla so that

$$(I \otimes \langle j|) U (|r\rangle_{\text{sys}} \otimes |0\rangle_{\text{anc}}) = K_j |r\rangle$$

for $|r\rangle$ in an orthonormal basis of $\text{supp}(\Sigma)$ (dimension r_Σ). We complete the isometry $V : \mathbb{C}^{r_\Sigma} \rightarrow \mathbb{C}^{(MN)^k \cdot N}$ to a full unitary via Householder QR of $[V \mid \text{random}]$ with sign-fixing so the first r_Σ columns match V exactly. This is then wrapped as a Qiskit `UnitaryGate` and applied inside an `AerSimulator` statevector circuit. An early implementation attempt used a rank-1 factorization and correctly returned empirical success far below the analytic value; see `report/failure_analysis.md` for the debugging history.

4 Results vs paper

4.1 Test A: analytic Eq. (15) sweep

See Figure ?? (figures/fig1_pgm_success_vs_N.png). Selected rows:

k	N	M	Pr(success)	Pr($1 \leq \eta_w^x \leq 4$)
2	32	5	0.5284	0.5266
2	64	8	0.6110	0.6151
3	27	3	0.6380	0.6508
3	32	3	0.5470	0.5579
4	16	2	0.6144	0.6302
4	24	2	0.4792	0.4889

The paper’s Lemma 2 predicts $\Pr(1 \leq \eta_w^x \leq 4)$ is a constant when $M = \lfloor N^{1/k} \rfloor$. For $k = 3$, whenever N is a perfect k -th power (so $M = \lfloor N^{1/k} \rfloor$ is exact), the fraction stays ≥ 0.55 ; when M under-shoots (e.g. $N = 32, k = 3, M = 3$ but $32^{1/3} \approx 3.17$), the fraction is a bit lower but still substantial. Our finite- N data trend supports the asymptotic constant lower bound.

4.2 Test B: Operational $\text{Tr}(E_s \sigma_s)$ vs analytic Eq. (15)

N	M	k	Hilbert dim	Eq. 15	$\text{Tr}(E_s \sigma_s)$	diff
3	2	2	36	0.765069	0.765069	3.3e-16
4	2	2	64	0.666513	0.666513	4.4e-16
3	2	3	216	0.886031	0.886031	6.7e-16
4	2	3	512	0.827449	0.827449	4.4e-16

This is the **main proof of the algorithm’s core correctness**: building the PGM POVM by the paper’s own definition and applying it to the ensemble states yields, to machine precision, the closed-form success probability of Eq. (15).

4.3 Test C: Qiskit end-to-end

N	M	k	s	shots	Hilbert dim	qubits	Eq. 15	empirical	emp-eq15
3	2	2	1	500	36	8	0.765	0.754	0.011
4	2	2	3	500	64	8	0.667	0.624	0.043
3	2	3	2	300	216	10	0.886	0.880	0.006

All three empirical values match Eq. (15) within the expected Monte Carlo standard error $\sqrt{p(1-p)/\text{shots}} \approx 0.02\text{--}0.03$. Qiskit’s own transpiler + **AerSimulator** statevector backend executed the circuit; the isometry-completion tolerance was $< 3 \times 10^{-14}$ and the joint unitary was verified unitary to the same tolerance.

4.4 Figures

5 Verdict

REPLICATED. The paper’s headline quantitative claims (Eq. 15 and Lemma 2, together giving the constant lower bound on PGM success probability underpinning Theorem 3) are reproduced numerically both by direct enumeration and by an end-to-end Qiskit statevector simulation of the




Figure 1: PGM success probability from Eq. (15) at $M = \lfloor N^{1/k} \rfloor$. The $k = 3$ and $k = 4$ curves stay well away from zero when $N^{1/k}$ is close to an integer (marked kinks at perfect powers); the $k = 2$ (dihedral) curve behaves similarly but corresponds to the abelian regime because $M = \lfloor N^{1/2} \rfloor$ is not the “hard” $M=2$ regime.

Naimark-dilated PGM, all matching to Monte-Carlo tolerance or machine precision (as appropriate). The Lenstra integer-programming subroutine (Section 4 of the paper) is not *re-implemented*, but its role in the algorithm is fulfilled here by brute-force enumeration of matrix-sum solutions (which is polynomially equivalent for the constant k we test) — so the algorithm’s asymptotic *complexity* is a claim we cannot verify at $N \leq 64$, but its *success probability* at those sizes is faithfully reproduced.

6 Open Questions

See `report/open_questions.json` for the machine-readable version.

[label=Q6.]

1. In our $N = 32, k = 3, M = 3$ run, $\Pr(1 \leq \eta_w^x \leq 4) = 0.558$, higher than the $\Pr(\text{success}) = 0.547$. Why the (small but systematic) gap? Lemma 1’s lower bound $\alpha\beta^2 N/M^k$ with $\alpha = 1, \beta = 0.558$ gives $\geq 0.558^2 \cdot 32/27 \approx 0.369$, well below the actual 0.547. *Basis*: the Lemma 1 bound is loose; the tightness gap grows with k , and a sharper analytic bound (perhaps involving $E[\sqrt{\eta_w^x}]^2$) would predict success within a few percent instead of a factor of ~ 1.5 . *Follow-up*: derive a Cauchy–Schwarz-free bound that uses the full η distribution.
2. For fixed $M = 2$ and $k = 4$, our data shows $\Pr(\text{success})$ decays roughly as $1/N^{0.85}$ over $N \in [8, 24]$, not the Lemma 1 upper bound $M^k/N = 16/N$ (which is only saturated asymptotically). At what N does the empirical curve cross the $16/N$ upper bound? Direct enu-




Figure 2: Lemma 2 empirical: $\Pr(1 \leq \eta_w^x \leq 4)$ vs N . Bounded below by a constant ≥ 0.4 at $k \geq 3$ across the tested range, consistent with the paper’s asymptotic bound.

meration is limited to $N \lesssim 40$ at $k = 4$; is there a way to extrapolate reliably from small N ?
Basis: Test A data for $k = 4, M = 2$.

3. The paper’s Naimark implementation of the PGM (Eqs. 21–27) assumes efficient “quantum sampling” from matrix-sum solutions. In our simulation we replace this by classical enumeration (Lenstra unused), then build a full unitary via random QR completion. If the quantum sampling step is *approximate* with fidelity $1 - \delta$, how does the PGM success probability degrade? A concrete perturbation experiment (add unitary noise of strength ϵ to the Kraus operators and measure success drop) would give a sensitivity curve the paper does not compute. *Basis:* the paper only proves the exact-implementation bound.
4. In Test C, our $(N, M, k) = (4, 2, 2)$ run had the largest empirical vs analytic gap (0.043). Is this just Monte Carlo noise ($\sigma \approx 0.02$ at 500 shots gives $\sim 5\%$ chance of a 0.04 deviation) or is there a bias from the specific choice of $s = 3$? *Basis:* the paper claims $\Pr(\text{success})$ is independent of s , but the operational Kraus operators K_j are constructed once and then applied to a specific ρ_s , so numerical asymmetries in the completion of the Naimark unitary could induce s -dependence. A systematic sweep over all $s \in \mathbb{Z}_N$ would settle it.
5. Extending to the “generalized graph isomorphism” problem suggested at the end of Section 5 of the paper: given $G_0, \dots, G_{M-1}$ related by $G_b = \pi(G_{b+1})$, can our Qiskit statevector prototype be adapted (with a graph-encoded oracle for $f(b, \cdot)$) to demonstrate the PGM on random small instances (e.g. $n = 4$ vertices, $M = 8$)? The paper leaves this as motivation but does not carry it out. *Basis:* the oracle-agnostic PGM machinery we built (Test C) is already general enough; only the state-preparation stage would need to change.




Figure 3: Success probability if M is held *below* $N^{1/k}$ (here $M = 2$ fixed). Success decays approximately as M^k/N (Lemma 1 upper bound), consistent with the paper's remark that the interesting open regime is $2 \leq M \leq N^\epsilon$.